
NEURAL NETWORKS

目录

| | |
|--------------------------------|----------|
| 1 Bases | 2 |
| 1.1 Architectures | 2 |
| 1.2 Loss functions | 3 |
| 1.3 Gradient Descent | 3 |
| 1.4 Backpropagation | 4 |

1 Bases

1.1 Architectures




图 1: neural-networks-architectures

We will focus on one node, by noting:

- i : the i -th layer of the network
- j : the j -th hidden units of the layer
- (\vec{x}, y) : datasets, where \vec{x} is the input and y is the desired output. $\vec{x} \in \mathcal{R}^{n_x}$ has n_x variables
- $\vec{w} \in \mathcal{R}^{n_x}$: weight, each w corresponds to one x
- $b \in \mathcal{R}$: bias
- z : output

We have:

Forward propagation (before using the activation function):

$$z_j^{[i]} = (\vec{w}_j^{[i]})^T \cdot \vec{x} + b_j^{[i]}$$

And then, **activation functions** $\hat{y} = g(z)$ are used at the end of a hidden unit to introduce non-linear complexities to the model.

| Sigmoid | Tanh | ReLU | Leaky ReLU |
|-------------------------------|--|---------------------|---|
| $g(z) = \frac{1}{1 + e^{-z}}$ | $g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$ | $g(z) = \max(0, z)$ | $g(z) = \max(\epsilon z, z)$ with $\epsilon \ll 1$ |
| | | | |

图 2: activation-functions

Suming up:

We have dataset (\vec{x}, y) . With input $x \in \mathcal{R}^{n_x}$, and the help of \vec{w}, b and g ,

$$\vec{x} \mapsto z \mapsto \hat{y} = g(z)$$

In this way having \hat{y} in our model and y in reality.

1.2 Loss functions

Here, we are going to measure the difference between y and \hat{y} . **Loss functions** are to measure how well our algorithm is doing based on one single sample.

Cross-entropy loss:

$$L(\hat{y}, y) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$$

Mean squared error loss:

$$L(\hat{y}, y) = (y - \hat{y})^2$$

Also, we have the **cost function** based on a number of samples:

Cost function: to measure how well you're doing in the entire training set. Here, n means the n -th sample, and m is the total number of the samples.

$$J(w, b) = \frac{1}{m} \sum_{n=1}^m L(\hat{y}^{[n]}, y^{[n]})$$

1.3 Gradient Descent

Gradient Descent is based on a convex function and tweaks its parameters iteratively to minimize a given function (here, the cost function) to its local minimum.




图 3: gradient-descent-1D




图 4: gradient-descent-3D

We note :

- a : current position

- b : the next position
- α : a waiting factor
- $\nabla f(a)$: the direction of the steepest descent at a

What **gradient descent** does:

$$b = a - \alpha \cdot \nabla f(a)$$

We define α as the **learning rate**, which indicates at which space the weights get updated.

It is very important for us to select a appropriate value of α .




图 5: gradient-descent-learning-rate

After several iterations by repeating the method until $|\nabla f| < \varepsilon$, which means it has converged, then stop.

1.4 Backpropagation

Backpropagation is an algorithm for supervised learning of artificial neural networks using **gradient descent**.

Here, we have :

$$\hat{y} = g(z), z = w^T x + b$$

For one single node with $(x \in \mathcal{R}^{n_x}, y)$, we focus on one variable w_k which correspond to $x_k \in \{x_1, \dots, x_{n_x}\}$ by using the **chain rule of differentiation**:

$$\begin{aligned} \nabla_w &= \frac{\partial L(\hat{y}, y)}{\partial w_k} = \frac{\partial L(\hat{y}, y)}{\partial \hat{y}} \cdot \frac{d\hat{y}}{dz} \cdot \frac{\partial z}{\partial w_k} \\ &= \frac{\partial L(\hat{y}, y)}{\partial z} \cdot x_k \end{aligned}$$

If g is the sigmoid function, and the loss function is the cross-entropy loss function, then for each w_k ,

$$\begin{aligned} \nabla_w &= \frac{\partial L(\hat{y}, y)}{\partial w_k} = (\hat{y} - y) \cdot x_k \\ \begin{cases} w_k & := w_k - \alpha(\hat{y} - y)x_k \\ b_k & := b_k - \alpha(\hat{y} - y) \end{cases} \end{aligned}$$

w and b are updated as follows :

1. Take a batch of training data
2. Perform **forward propagation** to obtain the corresponding loss
3. **Backpropagate** the loss to get the gradients
4. Use the gradients to update the weights of the network

The whole process is:

```
---Initializing---
g = (an activation function)
J = 0
L = (Cross-entropy loss)
For p = 1 to (the total number of inputs)
    dw_p = 0
    w_p = (random but appropriate number)
db = 0
alpha = (learning rate)

---For all the samples---
For n = 1 to m
    ---Forward Propagation---
    z = wx + b
    a = g(z)
    J += L(a,y)
    ---Backpropagation---
    dz = (after calculation)
    For p = 1 to (the total number of inputs)
        dw_p += x_p * dz
    db += dz

---Update w and b by using the gradients---
J /= m
For p = 1 to (the total number of inputs)
    dw_p /= m
    w_p := w_p - alpha * dw_p
b := b - alpha * db
```